

# Homework 3

School & Student ID: NYCU / 513 559 004  
Course Title: Machine Learning  
Announced Date: 241007  
Deadline: 241021

---

## Table of Contents

<b>Homework 3.....</b>	<b>1</b>
Question 5.....	2
Question 6.....	3
Question 7.....	4
Question 8.....	5
Question 9.....	7
Question 10.....	8
Question 11.....	10
Question 12.....	12
Question 13.....	14

---

---

## Question 5

Suppose the union  $H_1 \cup H_2$  can leverage hypotheses from both sets to shatter more points than either set could alone.

$$d_{vc}(H_1 \cup H_2) > d_{vc}(H_1) + d_{vc}(H_2)$$

However, for any labeling on a set of  $d_{vc}(H_1) + d_{vc}(H_2)$  points, at least one of the hypothesis sets  $H_1$  or  $H_2$  must fully accommodate that labeling.

Thus, the inequality above would require one of the sets to shatter more points than its individual VC dimension, which is impossible.

Therefore,

$$d_{vc}(H_1 \cup H_2) \leq d_{vc}(H_1) + d_{vc}(H_2)$$

---

## Question 6

Given that,

$$f_{0/1}(x) = \text{sign}\left(P(y = +1 | x) - \frac{1}{2}\right)$$

It implies that we classify  $y = +1$  only if  $P(y = +1 | x) > 0.5$ .

Then, we adjust the classification rule with asymmetric error costs:

$$f_{MKT}(x) = \text{sign}(P(y = +1 | x) - \alpha)$$

Since false negatives are 10 times more costly, it is necessary to have a threshold much more likely to classify as +1 to avoid the penalty:

$$\begin{aligned}\frac{P(y=+1 | x)}{P(y=-1 | x)} &\geq 10 \\ \Rightarrow \frac{P(y=+1 | x)}{1-P(y=+1 | x)} &\geq 10 \\ \Rightarrow P(y = +1 | x) &\geq \frac{10}{11}\end{aligned}$$

Hence, we can put  $P(y = +1 | x)$  into the function  $f_{MKT}(x)$ , and we'll find:

$$\alpha = \frac{10}{11}$$

---

## Question 7

Given that,

$$E_{out}^{(2)}(h) = E_{x \sim P(x), y \sim (y|x)}[h(x) \neq y]$$
$$\Rightarrow E_{out}^{(2)}(h) = E_{x \sim P(x)}[P(h(x) \neq y | x)]$$

Decomposing it with the law of total probability:

$$P(h(x) \neq y | x) = P(h(x) \neq f(x) \text{ and } f(x) = y | x) + P(f(x) \neq y | x)$$

By linearity of expectation:

$$E_{out}^{(2)}(h) = E_x[P(h(x) \neq f(x) \text{ and } f(x) = y | x)] + E_x[P(f(x) \neq y | x)]$$

The first term relates to  $E_{out}^{(1)}(h)$ , as:

$$P(h(x) \neq f(x) \text{ and } f(x) = y | x) \leq P(h(x) \neq f(x) | x)$$

$$\Rightarrow E_x[P(h(x) \neq f(x) \text{ and } f(x) = y | x)] \leq E_{out}^{(1)}(h)$$

Combining components, therefore:

$$E_{out}^{(2)}(h) \leq E_{out}^{(1)}(h) + E_{out}^{(2)}(f)$$

---

## Question 8

Knowing that,

$$W_{LIN} = (X^T X)^{-1} X^T y$$

The matrix  $X$  can be written as:

$$X = \begin{pmatrix} 1 & x_{11} & x_{12} & \dots & x_{1d} \\ 1 & x_{21} & x_{22} & \dots & x_{2d} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_{N1} & x_{N2} & \dots & x_{Nd} \end{pmatrix}$$

Using the diagonal matrix  $X' = XD$  to scale the intercept term by 1126.

$$D = \begin{pmatrix} 1126 & 0 & 0 & \dots & 0 \\ 0 & 1 & 0 & \dots & 0 \\ 0 & 0 & 1 & \dots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \dots & 1 \end{pmatrix}$$

$$X' = \begin{pmatrix} 1126 & x_{11} & x_{12} & \dots & x_{1d} \\ 1126 & x_{21} & x_{22} & \dots & x_{2d} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1126 & x_{N1} & x_{N2} & \dots & x_{Nd} \end{pmatrix}$$

After changing  $x_0$  to 1126, the weight vector  $W_{LUCKY}$  for  $X'$  is shown as:

$$W_{LUCKY} = \left( (X')^T X' \right)^{-1} (X')^T y$$

Since  $X' = XD$ , we can substitute  $X'$  with  $XD$  as follows:

$$\Rightarrow W_{LUCKY} = \left( D^T X^T X D \right)^{-1} \left( D^T X^T y \right)$$

We know that  $\left( D^T X^T X D \right)^{-1} = D^{-1} \left( X^T X \right)^{-1} D^{-1}$ , we can replace it into the equation for  $W_{LUCKY}$ :

$$W_{LUCKY} = D^{-1} \left( X^T X \right)^{-1} X^T y$$

$$\Rightarrow W_{LIN} = D \cdot W_{LUCKY}$$

---

## Question 9

Given that,

$$\hat{h}(x) = \frac{1}{2} \left( \frac{w^T x}{\sqrt{1 + (w^T x)^2}} + 1 \right)$$

Minimize the negative log-likelihood for the data  $\{(x_n, y_n)\}$  :

$$\hat{E}_{in}(w) = -\frac{1}{N} \left( y_n \ln(\hat{h}(x)) + (1 - y_n) \ln(1 - \hat{h}(x)) \right)$$

Let  $u = w^T x$ .

$$\frac{d\hat{h}}{du} = \frac{1}{2} \cdot \frac{(1+u^2)-u^2}{(1+u^2)^{\frac{3}{2}}} = \frac{1}{2} \cdot \frac{1}{(1+u^2)^{\frac{3}{2}}}$$

Then, the gradient  $\nabla_w \hat{h}$  we find out would be:

$$\Rightarrow \nabla_w \hat{h} = \frac{d\hat{h}}{du} \cdot x$$

We can calculate the gradient  $\nabla \hat{E}_{in}(w)$  :

$$\nabla \hat{E}_{in}(w) = -\frac{1}{N} \sum_{n=1}^N \left( \frac{y_n}{\hat{h}(x_n)} - \frac{1-y_n}{1-\hat{h}(x_n)} \right) \nabla_w \hat{h}(x_n)$$

Substitute  $\nabla_w \hat{h}(x)$  :

$$\Rightarrow \nabla \hat{E}_{in}(w) = -\frac{1}{N} \sum_{n=1}^N \left( \frac{y_n - \hat{h}(x_n)}{\hat{h}(x_n)(1-\hat{h}(x_n))} \right) \cdot \frac{1}{2} \cdot \frac{x_n}{(1+(w^T x_n)^2)^{\frac{3}{2}}}$$


---

## Question 10

The first page of the snapshot of my *code is on the next page*.

My findings:

1. The scatter plot shows the relationship between Ein and Eout for different experiments.
2. Generally, Ein and Eout are *positively correlated*, meaning that when the model performs well on training data, it also tends to perform well on unseen data.
3. However, some variability indicates that a low Ein does not always guarantee a low Eout. This suggests that *overfitting can still occur* in some cases, depending on the specific training samples.





```

1  import numpy as np
2  import matplotlib.pyplot as plt
3  from sklearn.linear_model import LinearRegression
4  from google.colab import drive
5
6  # Load the dataset
7  drive.mount('/d')
8  file_path = '/d/MyDrive/ML_hw3_attachment-cpusmall_scale.txt'
9  data = []
10 with open(file_path, 'r') as f:
11     for line in f:
12         values = line.strip().split()
13         y = float(values[0])
14         x = [float(v.split(':')[1]) for v in values[1:]]
15         data.append((x, y))
16
17 X_full = np.array([x for x, y in data])
18 y_full = np.array([y for x, y in data])
19
20 # Parameters
21 N = 32
22 num_experiments = 1126
23 Ein_list = []
24 Eout_list = []
25
26 # Run the experiments
27 for _ in range(num_experiments):
28     # Randomly sample N examples for training
29     indices = np.random.choice(len(X_full), N, replace=False)
30     X_train = X_full[indices]
31     y_train = y_full[indices]
32
33     # Add bias term ( $x_0 = 1$ )
34     X_train = np.hstack((np.ones((X_train.shape[0], 1)), X_train))
35     X_full_bias = np.hstack((np.ones((X_full.shape[0], 1)), X_full))
36
37     # Fit linear regression model
38     model = LinearRegression(fit_intercept=False)
39     model.fit(X_train, y_train)
40     w_lin = model.coef_
41
42     # Calculate Ein (in-sample error)
43     y_train_pred = X_train @ w_lin
44     Ein = np.mean((y_train - y_train_pred) ** 2)
45     Ein_list.append(Ein)
46

```

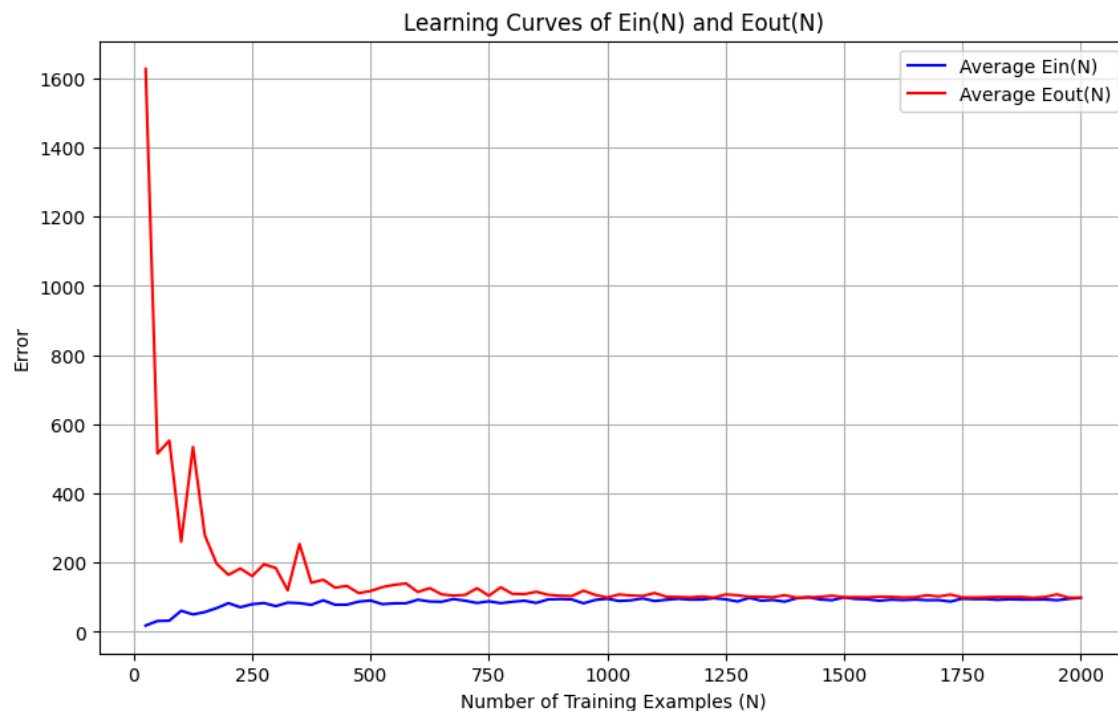
---

## Question 11

The first page of the snapshot of my *code is on the next page*.

My findings:

1. As  $N$  increases,
  - a. both  $E_{in}$  and  $E_{out}$  tend to **decrease**, indicating improved model generalization with more training data.
  - b. The gap between  $E_{in}$  and  $E_{out}$  **narrows**, suggesting reduced overfitting with larger training sizes.
  - c. Both **converge** to similar values, indicating that the model is better at generalizing to new data, reducing bias and variance.
2. For smaller values of  $N$ , there is a significant difference between  $E_{in}$  and  $E_{out}$ , which indicates **overfitting**. *The model fits the training data well but struggles to generalize to unseen data.*
3. The learning curve shows that **adding more training data** helps the model improve its performance, especially for smaller training set sizes where the model initially suffers from high variance.



```

10     with open(file_path, 'r') as f:
11         for line in f:
12             values = line.strip().split()
13             y = float(values[0])
14             x = [float(v.split(':')[1]) for v in values[1:]]
15             data.append((x, y))
16
17     X_full = np.array([x for x, y in data])
18     y_full = np.array([y for x, y in data])
19
20     # Parameters
21     N_values = np.arange(25, 2001, 25)
22     num_experiments = 16
23     average_Ein = []
24     average_Eout = []
25
26     # Run experiments for each value of N
27     for N in N_values:
28         Ein_total = 0
29         Eout_total = 0
30
31         for _ in range(num_experiments):
32             # Randomly sample N examples for training
33             indices = np.random.choice(len(X_full), N, replace=False)
34             X_train = X_full[indices]
35             y_train = y_full[indices]
36
37             # Add bias term ( $x_0 = 1$ )
38             X_train = np.hstack((np.ones((X_train.shape[0], 1)), X_train))
39             X_full_bias = np.hstack((np.ones((X_full.shape[0], 1)), X_full))
40
41             # Fit linear regression model
42             model = LinearRegression(fit_intercept=False)
43             model.fit(X_train, y_train)
44             w_lin = model.coef_
45
46             # Calculate Ein (in-sample error)
47             y_train_pred = X_train @ w_lin
48             Ein = np.mean((y_train - y_train_pred) ** 2)
49             Ein_total += Ein
50
51             # Calculate Eout (out-of-sample error)
52             y_full_pred = X_full_bias @ w_lin
53             Eout = np.mean((y_full - y_full_pred) ** 2)

```


---

## Question 12

The first page of the snapshot of my *code is on the next page*.

My findings:

1. With only the first 2 features, both  $E_{in}$  and  $E_{out}$  are ***generally higher compared to using all 12 features***, indicating a decrease in model performance.
2. The gap between  $E_{in}$  and  $E_{out}$  ***remains more significant for smaller values of  $N$*** , suggesting that the reduced feature set leads to increased bias and potentially higher variance.
3. As  $N$  increases, the errors decrease, but they ***do not reach as low values as when using all features***, demonstrating that the model benefits from having more features to learn from.



```

17     X_full = np.array([x[:2] for x, y in data]) # Use only the first 2 features
18     y_full = np.array([y for x, y in data])
19
20     # Parameters
21     N_values = np.arange(25, 2001, 25)
22     num_experiments = 16
23     average_Ein = []
24     average_Eout = []
25
26     # Run experiments for each value of N
27     for N in N_values:
28         Ein_total = 0
29         Eout_total = 0
30
31         for _ in range(num_experiments):
32             # Randomly sample N examples for training
33             indices = np.random.choice(len(X_full), N, replace=False)
34             X_train = X_full[indices]
35             y_train = y_full[indices]
36
37             # Add bias term ( $x_0 = 1$ )
38             X_train = np.hstack((np.ones((X_train.shape[0], 1)), X_train))
39             X_full_bias = np.hstack((np.ones((X_full.shape[0], 1)), X_full))
40
41             # Fit linear regression model
42             model = LinearRegression(fit_intercept=False)
43             model.fit(X_train, y_train)
44             w_lin = model.coef_
45
46             # Calculate Ein (in-sample error)
47             y_train_pred = X_train @ w_lin
48             Ein = np.mean((y_train - y_train_pred) ** 2)
49             Ein_total += Ein
50
51             # Calculate Eout (out-of-sample error)
52             y_full_pred = X_full_bias @ w_lin
53             Eout = np.mean((y_full - y_full_pred) ** 2)
54             Eout_total += Eout
55
56         # Calculate average Ein and Eout
57         average_Ein.append(Ein_total / num_experiments)
58         average_Eout.append(Eout_total / num_experiments)
59
60     # Plot learning curves

```

---

### Question 13

Knowing that  $B(N, k)$  covers every possible subset size from 0 up to  $k - 1$ .

$$B(N, k) \leq \sum_{i=0}^{k-1} C_i^N$$

Precisely,  $B(N, k)$  should count all subsets that can be formed up to  $k$  elements. Such that, it appears  $B(N, k)$  is at least large enough to include all such subsets.

$$B(N, k) = \sum_{i=0}^k C_i^N \geq \sum_{i=0}^{k-1} C_i^N$$

$$\Rightarrow B(N, k) \geq \sum_{i=0}^{k-1} C_i^N$$

Since  $B(N, k)$  includes all subsets up to that size, it cannot be less than  $\sum_{i=0}^{k-1} C_i^N$ . In

addition, we have already known  $B(N, k) \leq \sum_{i=0}^{k-1} C_i^N$ , we can eventually establish:

$$B(N, k) = \sum_{i=0}^{k-1} C_i^N$$